

# System Integration - Timed Automata


Hugo A. López

[hulo@dtu.dk](mailto:hulo@dtu.dk)


**A little recap**

# This course so far

+ Operational  
Easier to verify



+ Abstract  
Harder to verify



# This course so far

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



# This course so far

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



Automata

Computation

# This course so far

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



Automata

Workflow Models

Computation

Task distribution

# This course so far

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



Automata

Workflow Models

Interaction Models

Computation

Task distribution

Handovers

# This course so far

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



Automata

Workflow Models

Interaction Models

Requirement Models

Computation

Task distribution

Handovers

Strategy



# This course so far

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



Automata

Workflow Models

Interaction Models

Requirement Models

Enterprise  
Architecture  
Models

Computation

Task distribution

Handovers

Strategy

Technology  
Alignment

# This class

+ Operational  
Easier to verify

+ Abstract  
Harder to verify



Automata

Workflow Models

Interaction Models

**Real-time models**

Requirement Models

Enterprise  
Architecture  
Models

Computation

Task distribution

Handovers

**Temporal Dependencies**

Strategy

Technology  
Alignment



# Why real-time systems



# Real-time systems

- Parts of the socio-technical system are expected to finish within reliable time bounds
- Communication with microservices does not run forever, but under standard service level agreements
- Fault-tolerant systems need to model availability considerations under absences of responses (failures)

A system where correctness does not only depend on the logical order of events, but also on their **timing**




# How does time influence our specification?

- Evolution



# How does time influence our specification?

- Evolution




## Büchi Automata

- Infinite alphabet
- Initial and accepting states
- Accept execution if pass through accepting state infinitely many times


# How does time influence our specification?

- Evolution



## Büchi Automata

- Infinite alphabet
- Initial and accepting states
- Accept execution if pass through accepting state infinitely many times




## Büchi Timed Automata

- Büchi-accepting
- Real-valued variables: modelling clock
- Constraints on clock variables and resets


# How does time influence our specification?

- Evolution




## Büchi Automata

- Infinite alphabet
- Initial and accepting states
- Accept execution if pass through accepting state infinitely many times



## Büchi Timed Automata

- Büchi-accepting
- Real-valued variables: modelling clock
- Constraints on clock variables and resets



## Timed Safety Automata

- Clock variables
- Local invariant conditions
- Accept when invariant is satisfied



# Timed Automata



- A timed automaton is a tuple

$$(\mathcal{L}, l_0, C, \mathcal{A}, \mathcal{E}, \mathcal{I})$$

- where  $\mathcal{L}$  is a set of locations,
- $l_0 \in \mathcal{L}$  is the initial location,
- $C$  is the set of clocks,
- $\mathcal{A}$  is a set of actions, co-actions and the internal  $\tau$ -action,
- $\mathcal{E} \subseteq \mathcal{L} \times \mathcal{A} \times \mathcal{B}(C) \times 2^C \times \mathcal{L}$  is a set of edges between locations with an action, a guard and a set of clocks to be reset, and
- $\mathcal{I} : \mathcal{L} \rightarrow \mathcal{B}(C)$  assigns invariants to locations

# Semantics

**The operational Semantics of a timed automaton is:**

- If  $u, u + d \in I(l)$  and  $d \in \mathbb{R}^+$ ,  
then  $\langle l, u \rangle \xrightarrow{d} \langle l, u + d \rangle$   

- If  $l \xrightarrow{\tau, \alpha, r} l', u \in g, u' = [r \mapsto 0]u$  and  $u' \in I(l')$ ,  
then  $\langle l, u \rangle \xrightarrow{\alpha} \langle l', u' \rangle$   


• **Notation:**  $\langle l, u \rangle$  is a state


$\langle l, u \rangle \xrightarrow{\alpha} \langle l', u' \rangle$  is a transition

# Operational Semantics


- For a TA  $(\mathcal{L}, \ell_o, C, \mathcal{A}, \mathcal{E}, \mathcal{I})$ , its semantics is given in terms of an LTS  $\langle S, s_o, \rightarrow \rangle$ , where
  - $S \subseteq \mathcal{L} \times \mathcal{R}^C$  is the set of states,
  - $s_o = (\ell_o, u_o)$  is the initial state, and
  - $\rightarrow \subseteq S \times \{\mathcal{R}_{\geq 0} \cup \mathcal{A}\} \times S$  is the transition relation such that:
    - $(\ell, u) \xrightarrow{d} (\ell, u+d)$  if  $\forall d' : 0 \leq d' \leq d \Rightarrow u+d' \in \mathcal{I}(\ell)$ , and
    - $(\ell, u) \xrightarrow{a} (\ell', u')$  if there exists  $e = (\ell, a, g, r, \ell') \in \mathcal{E}$  such that
      - $u \in g$ ,
      - $u' = [r] \rightarrow 0 u$ , and
      - $u' \in \mathcal{I}(\ell')$

The transition respect timed invariants

# Our first CPS



Lamp





User



- System = parallel composition of lamp and user



# Our first CPS



- System = parallel composition of lamp and user

Question: how can we model the policy that the light should be dimmable if the user presses **quickly**?

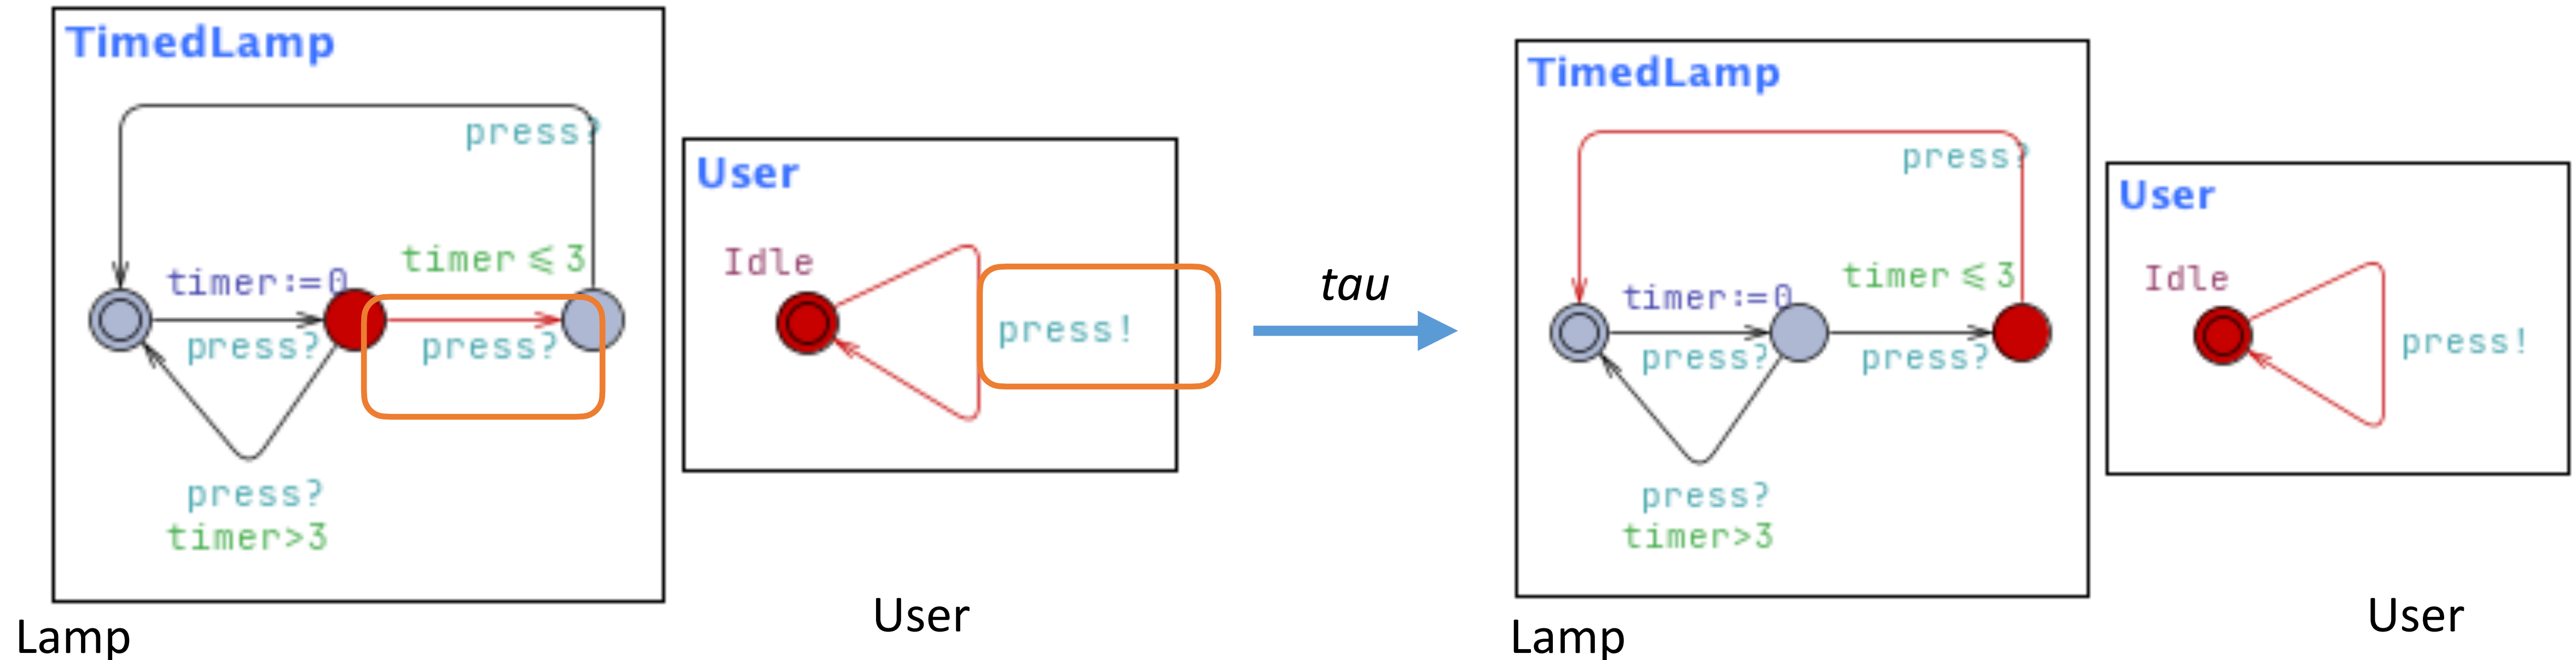
# Adding timers & guards



- A real-value timer measures the delay between the *press* events
- Clocks can be reset, and queried
- Multiple clocks can be instantiated



# Networks of Timed Automata

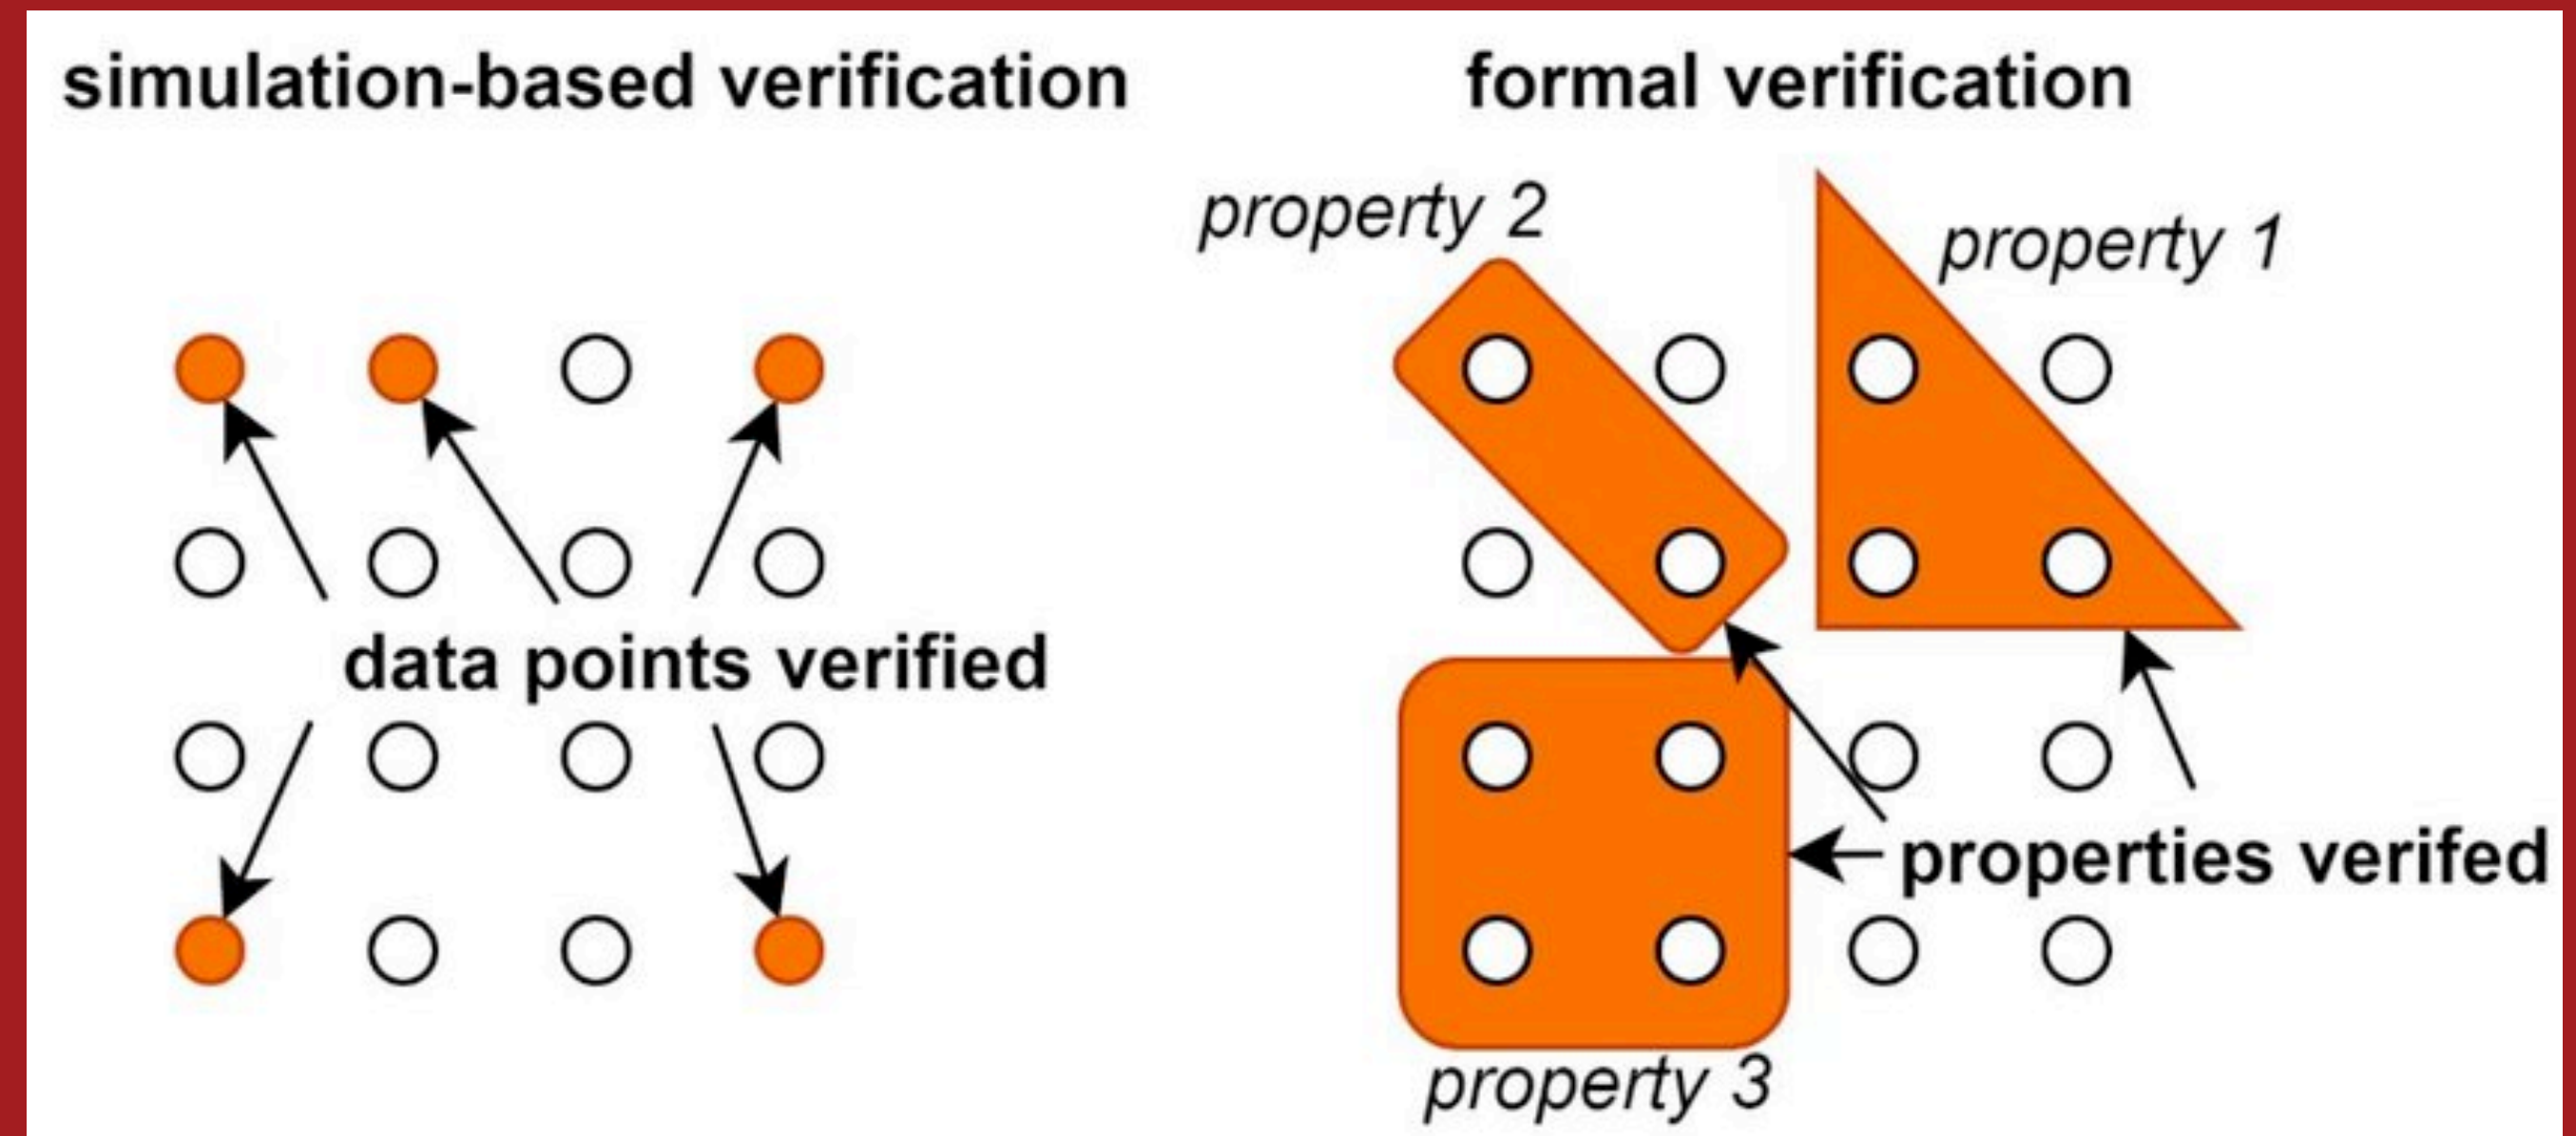



- A network of timed automata provides a two-way synchronization between co-actions
  - Akin to CCS!

# Extensions

- See tutorial for more details
  - Urgent channels: no delay if transition with urgent action can be taken.
  - Committed locations: reduce the number of clocks and disallow delays between committed locations
  - Broadcast channels: one-to-many communication
  - Parameterised templates: allow to span finite copies of a model





# Testing and Verification

A testing approach does not guarantee a fail-free implementation. It guarantees that the software passes the tests you have provided.

If you need full-assurance, you need higher guarantees, for instance, formal methods

# Verification of Timed Models

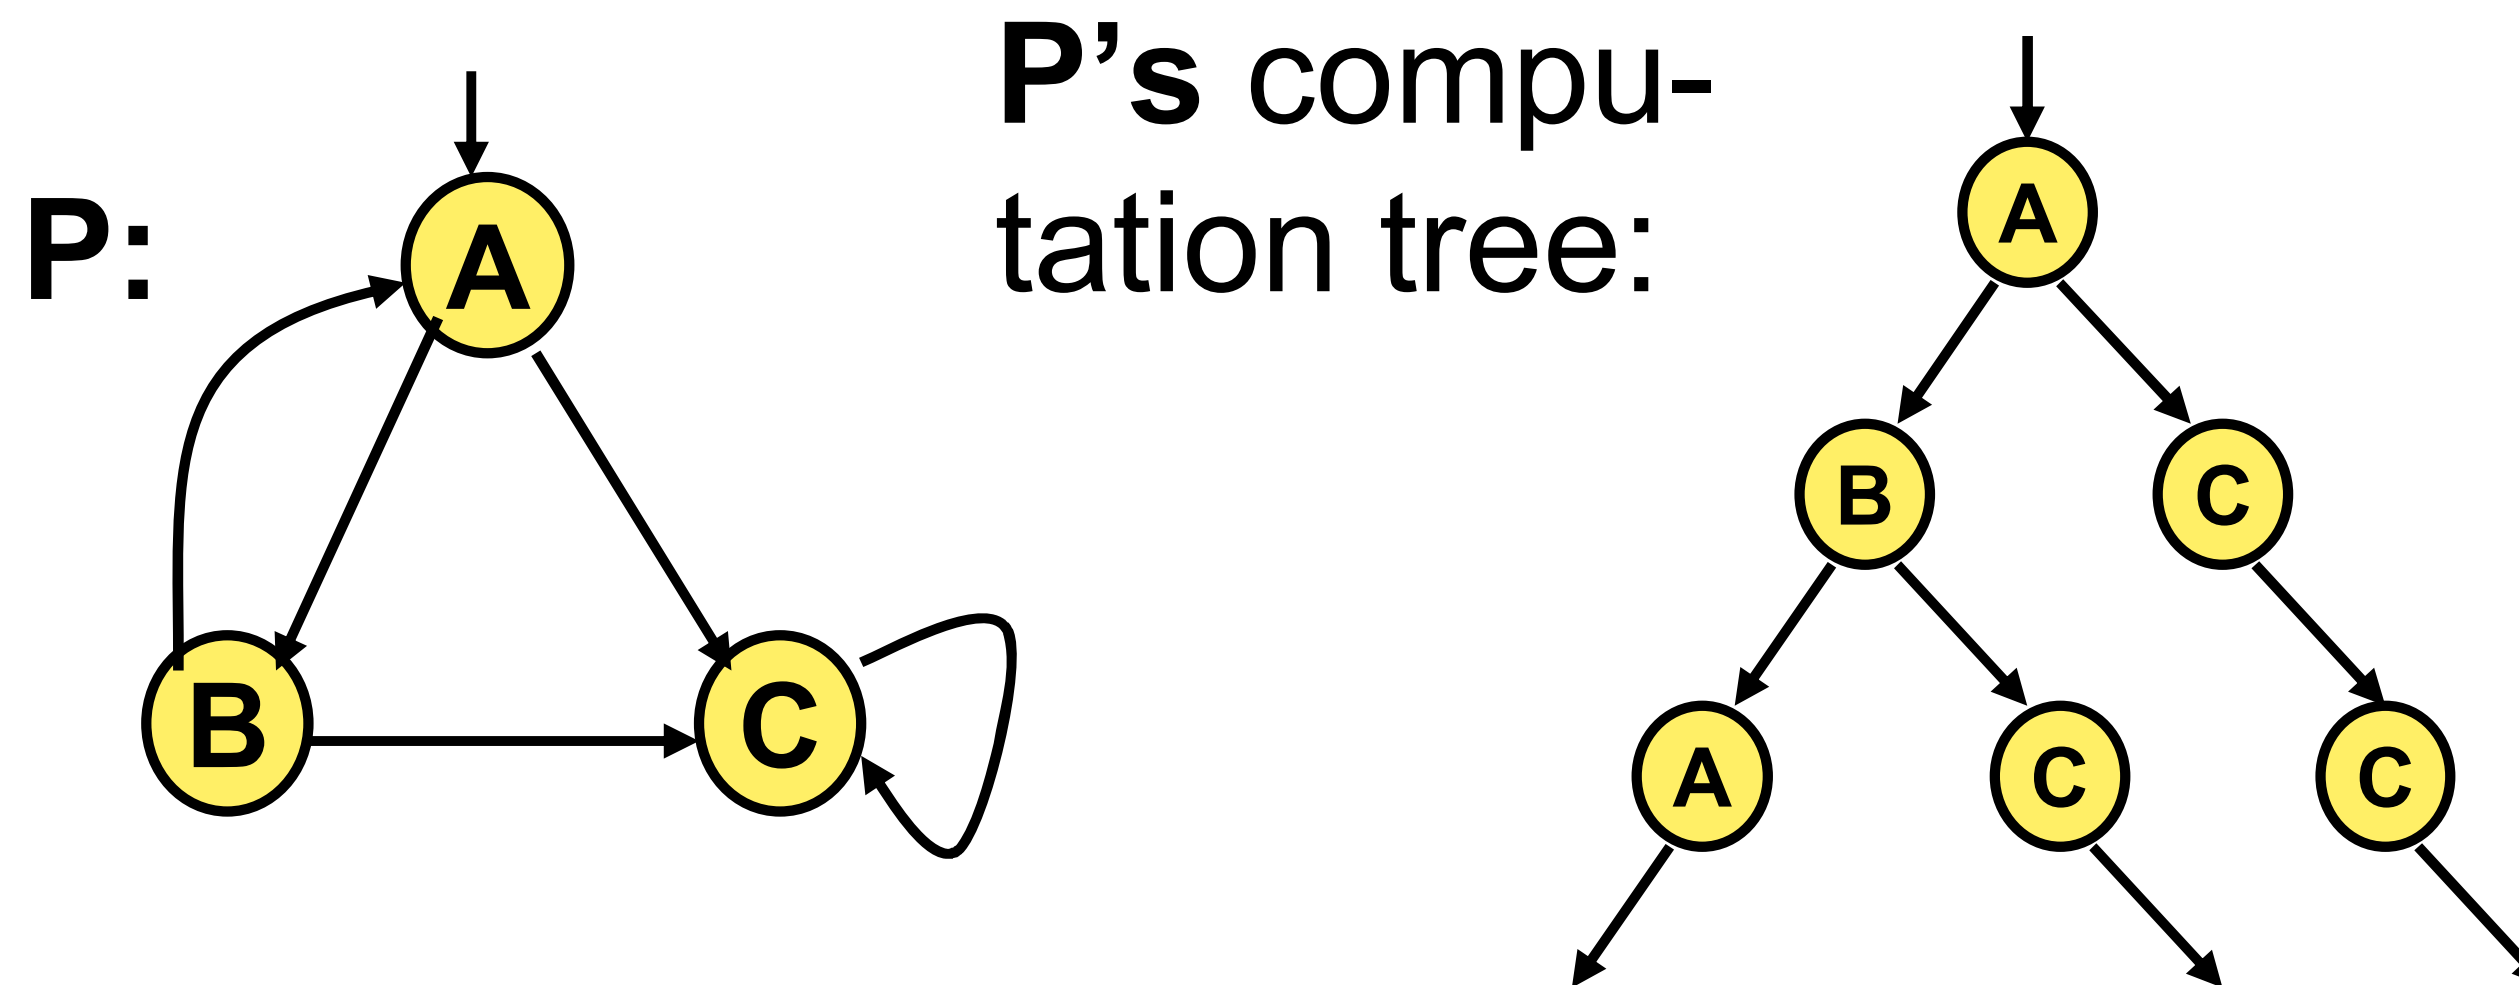
- Instead of running tests, we would like to apply **model-checking**
- Types of properties:
  - Safety: “**Something bad never happens**” [Lamport 1977]
    - Type 1: a broken invariant in the specification as state properties that fail
      - E.g. A mutual exclusion algorithm that allows more than 1 processors in the critical section
  - Every invariant is a safety property, but not the reverse
    - For e.g. "vending machine should get money **before** dispensing a product” is not a state property, but a **temporal property**

# Liveness

- Doing nothing easily fulfils a safety property as this will never lead to a “bad” situation
  - For e.g. "vending machine should get money **before** dispensing a product” not buying a product satisfy the property
- Liveness properties complement safety properties and require progress
  - “**something good**” will happen eventually [Lamport 1977]
- Liveness properties are violated in “infinite time”
  - whereas safety properties are violated in finite time
  - finite traces are of no use to decide whether  $P_{live}$  holds or not
  - any finite prefix can be extended such that the resulting infinite trace satisfies  $P_{live}$
  - Examples: fairness


# Verification in Uppaal

- A model checker verifies whether a model respects a requirement
- UPPAAL uses a simplified version of CTL [1] (temporal first-order logic)
- State formulae
- Path formulae: reachability, safety, liveness



# Formulae in TCTL


- Reachability  $\rightarrow E\langle\rangle P$
- Reads as “it is possible to reach a state in which  $P$  is satisfied



$P$  is true in at least one reachable state

# Formulae in TCTL


- Invariance:  $A[] P$
- “Always  $P$ ”



$P$  holds in each possible state

# Formulae in TCTL

- Inevitable  $P$ ,  $A \leftrightarrow P$
- $P$  will always eventually become true




$P$  is true in all paths, in some state



# Formulae in TCTL

- Potentially Always:  $E[] P$ .
- There exists a path in which  $P$  is true for each transition





# Formulae in TCTL

➤  $A[]p$ ,  $A\langle\rangle p$ ,  $E\langle\rangle p$ ,  $E[]p$  -  $p$  is a local property

➤ **Syntax:**




**automata location**      **data guard**      **clock guard**

$p ::= a.l \mid g_d \mid g_c \mid \text{deadlock} \mid$   
 $p \text{ and } p \mid p \text{ or } p \mid \text{not } p \mid$   
 $p \text{ imply } p \mid ( p )$


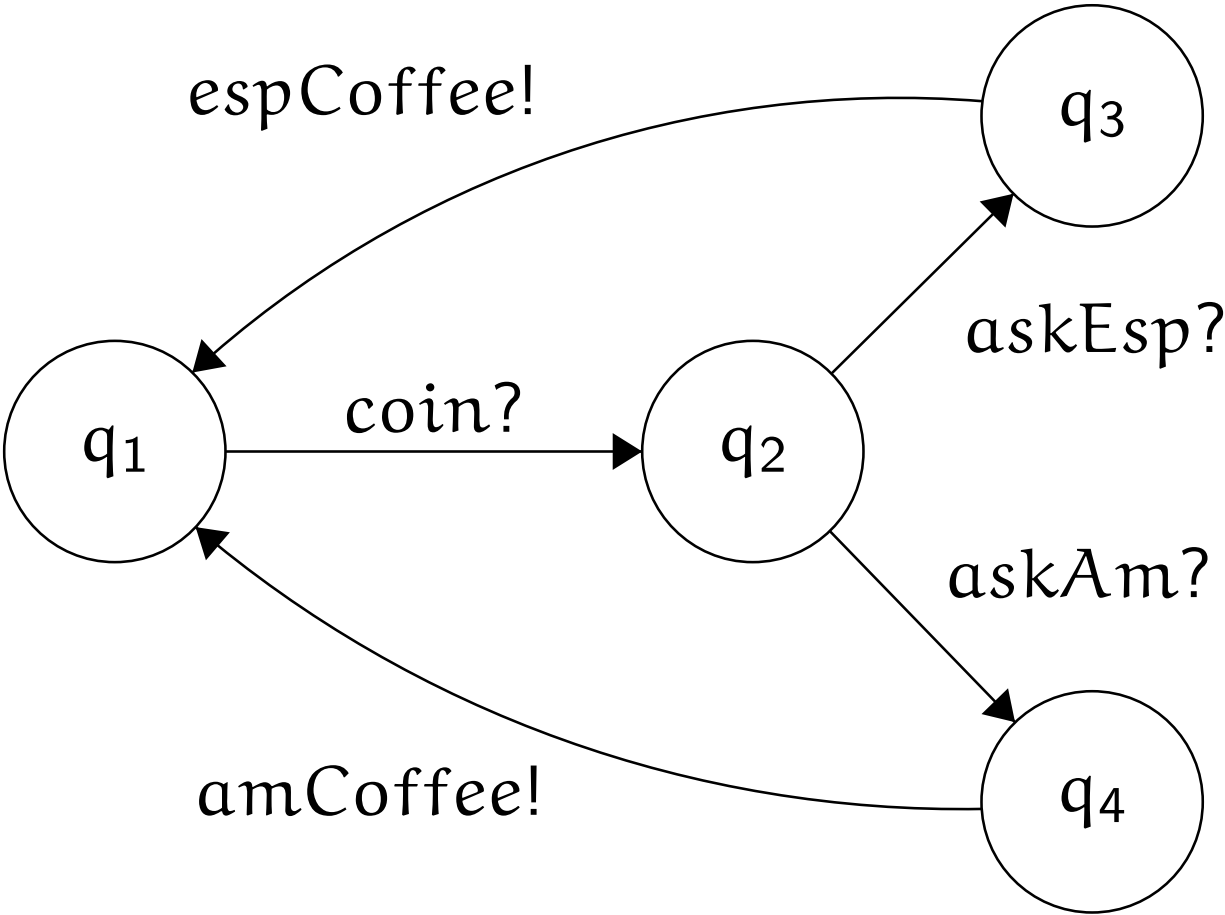
no action transition going  
out of a state or of its delay  
successors

**process name**

# Uppaal Demo: A lamp



# Second try: the coffee vending machine



# Take home message

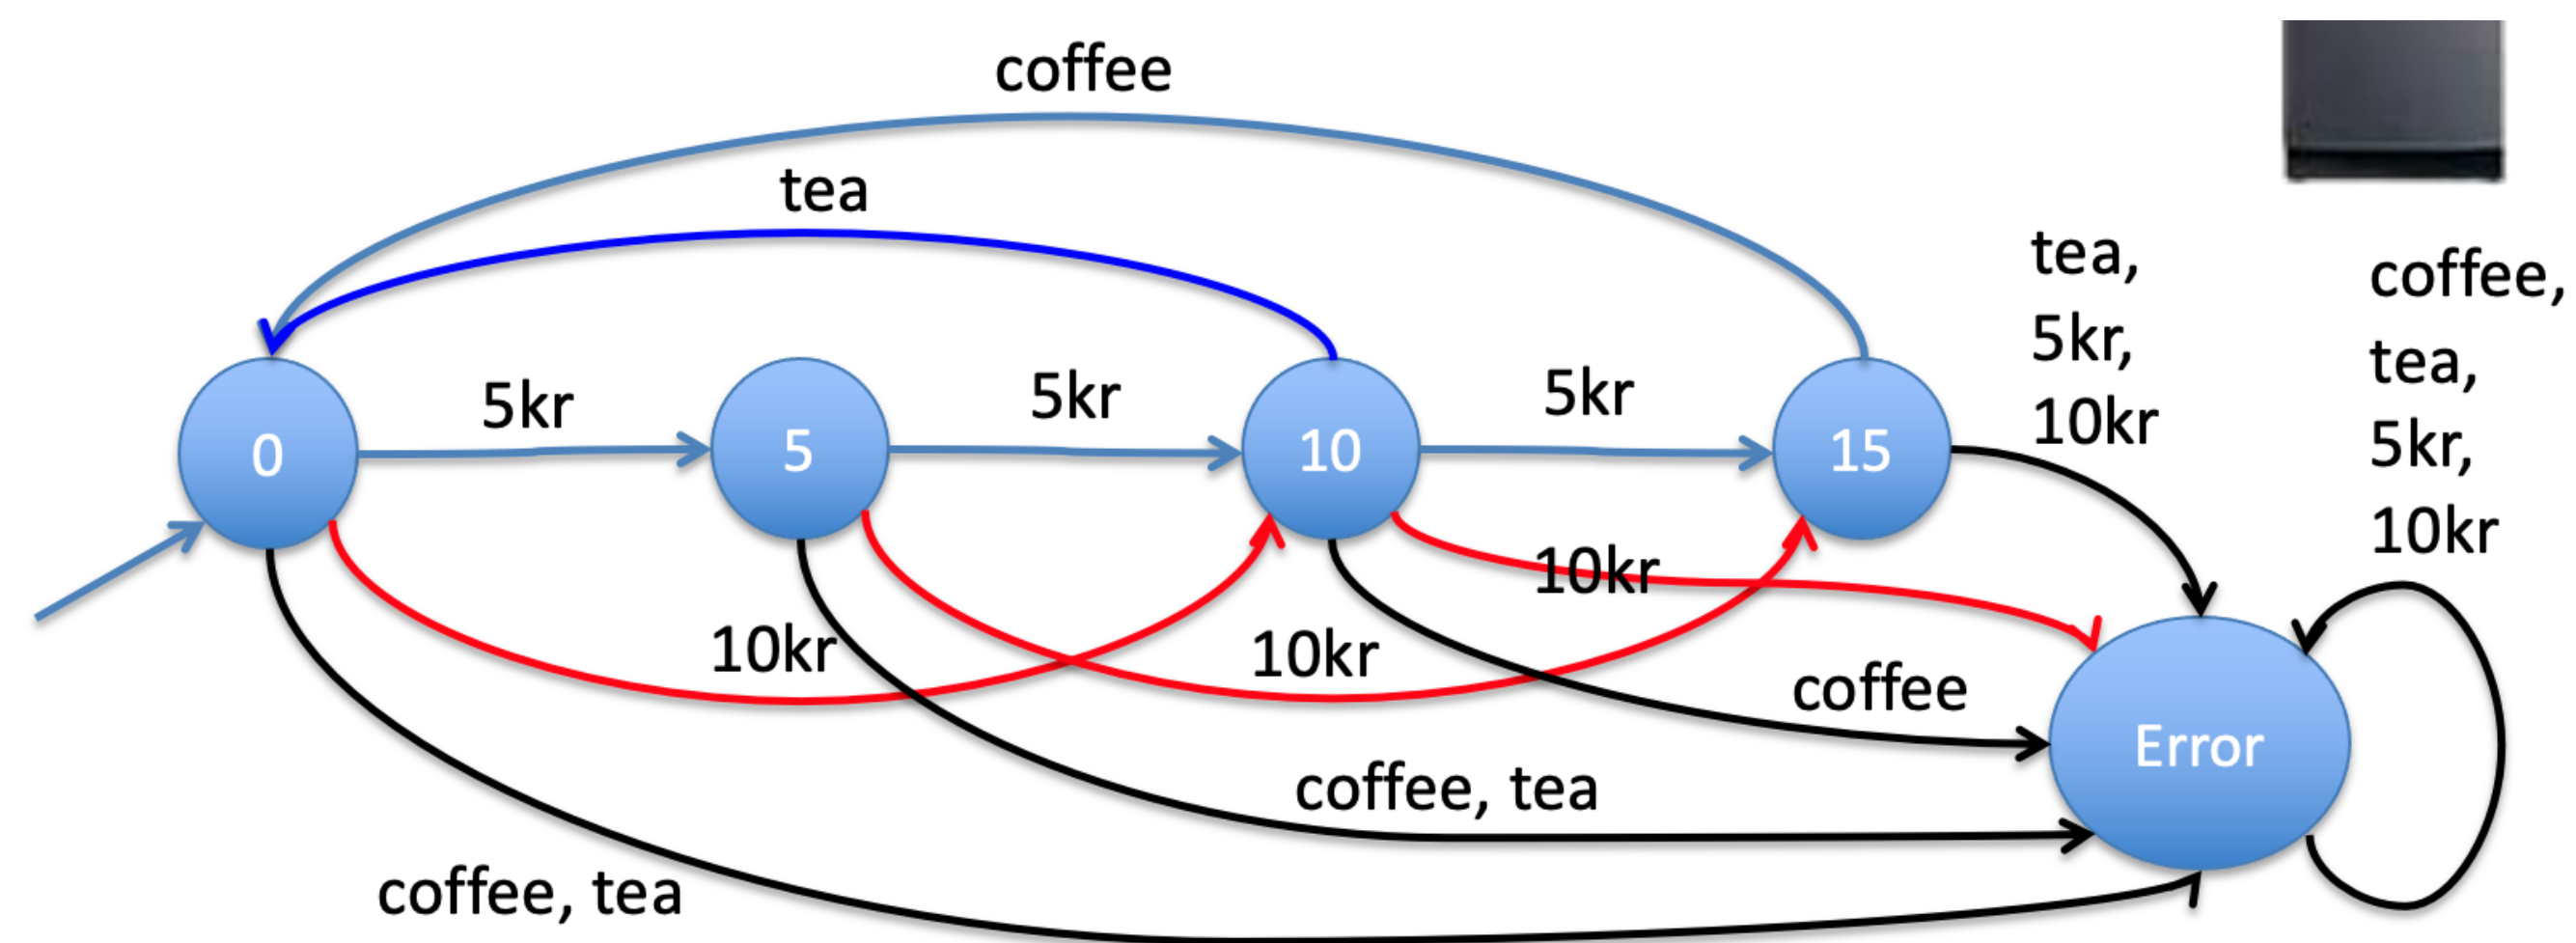
- Real time systems combine action-determinism, message passing, and time operations in to be considered correct
- While the state space becomes indeed humongous, we can verify the correctness of the system via model checking techniques

# Exercise

- Follow the tutorial "A Tutorial on Uppaal"  
Gerd Behrmann, Alexandre David, and Kim G. Larsen, and implement the models in sections 4 and 5.

# Exercise

- Implement the following state machine in uppaal, together with a corresponding set of users (that like tea or coffee)
- Make sure that the vending machine returns money after a timeout of 30 seconds
- Verify that for a multiple user can eventually be served by the vending machine (fairness)



# Exercise

- Consider the interaction model you handed in
  - Discuss with your group the type of timing constraints you may have in your model
    - And if you can, implement the model in Uppaal
  - Discuss with your group the type of model-checking properties you may be interested
    - And If you manage to implement the models, try to verify the new properties



# Peer review

- Carried out individually
- Step 1: each student submits a copy of the group project report and of the models they already submitted in the Assignment section.
- Submission through the Peer-review part 1 module in the Peer-review section
- All students working in the same group will submit the same files
- Make sure that the report and models are anonymized
- Deadline: March 14 at 11:59 (A.M.)



# Peer review

- Step 2: each student reviews 1 report of another group.
- Review through the Peer-review part 1 module in the Peer-review section
- All students working in the same group will get different reports
- Assessment is guided: a rubric of several criteria is defined
- You need to assign a score for each criterion
- You need to motivate your score, pointing out where issues are present in the models
- Feedback is anonymous to students (but not to teachers)
- Deadline: March 19 at 23:59.
- If you get more than 1 report to review, please contact the teachers.